

## Add destination

### Destination configuration

Send invocation records to a destination when your function is invoked asynchronously, or if your function processes records from a stream.

#### Source

The type of invocation that maps to the destination.

- ☒ Asynchronous invocation
- ☐ Stream invocation

#### Condition

The condition for using the destination.

- ☒ On failure
- ☐ On success

#### Destination type

An SQS queue, SNS topic, Lambda function, or EventBridge event bus.

SNS topic ▼

#### Destination

▼



Cancel

Save

## Add destination

### Destination configuration

Send invocation records to a destination when your function is invoked asynchronously, or if your function processes records from a stream.

#### Source

The type of invocation that maps to the destination.

- ☒ Asynchronous invocation  
☐ Stream invocation

#### Condition

The condition for using the destination.

- ☒ On failure  
☐ On success

#### Destination type

An SQS queue, SNS topic, Lambda function, or EventBridge event bus.

SNS topic ▼

#### Destination

▼



SNS, SQS, Lambda,  
EventBridge

Cancel

Save

```
{
  "version": "1.0",
  "timestamp": "2019-11-24T21:52:47.333Z",
  "requestContext": {
    "requestId": "8ea123e4-1db7-4aca-ad10-d9ca1234c1fd",
    "functionArn": "arn:aws:lambda:sa-east-1:123456678912:function:event-destinations:$LATEST",
    "condition": "RetriesExhausted",
    "approximateInvokeCount": 3
  },
  "requestPayload": {
    "Success": false
  },
  "responseContext": {
    "statusCode": 200,
    "executedVersion": "$LATEST",
    "functionError": "Handled"
  },
  "responsePayload": {
    "errorMessage": "Failure from event, Success = false, I am failing!",
    "errorType": "Error",
    "stackTrace": [
      "exports.handler (/var/task/index.js:18:18)"
    ]
  }
}
```

```
{
  "version": "1.0",
  "timestamp": "2019-11-24T21:52:47.333Z",
  "requestContext": {
    "requestId": "8ea123e4-1db7-4aca-ad10-d9ca1234c1fd",
    "functionArn": "arn:aws:lambda:sa-east-1:123456678912:function:event-destinations:$LATEST",
    "condition": "RetriesExhausted",
    "approximateInvokeCount": 3
  },
  "requestPayload": {
    "Success": false
  },
  "responseContext": {
    "statusCode": 200,
    "executedVersion": "$LATEST",
    "functionError": "Handled"
  },
  "responsePayload": {
    "errorMessage": "Failure from event, Success = false, I am failing!",
    "errorType": "Error",
    "stackTrace": [
      "exports.handler (/var/task/index.js:18:18)"
    ]
  }
}
```



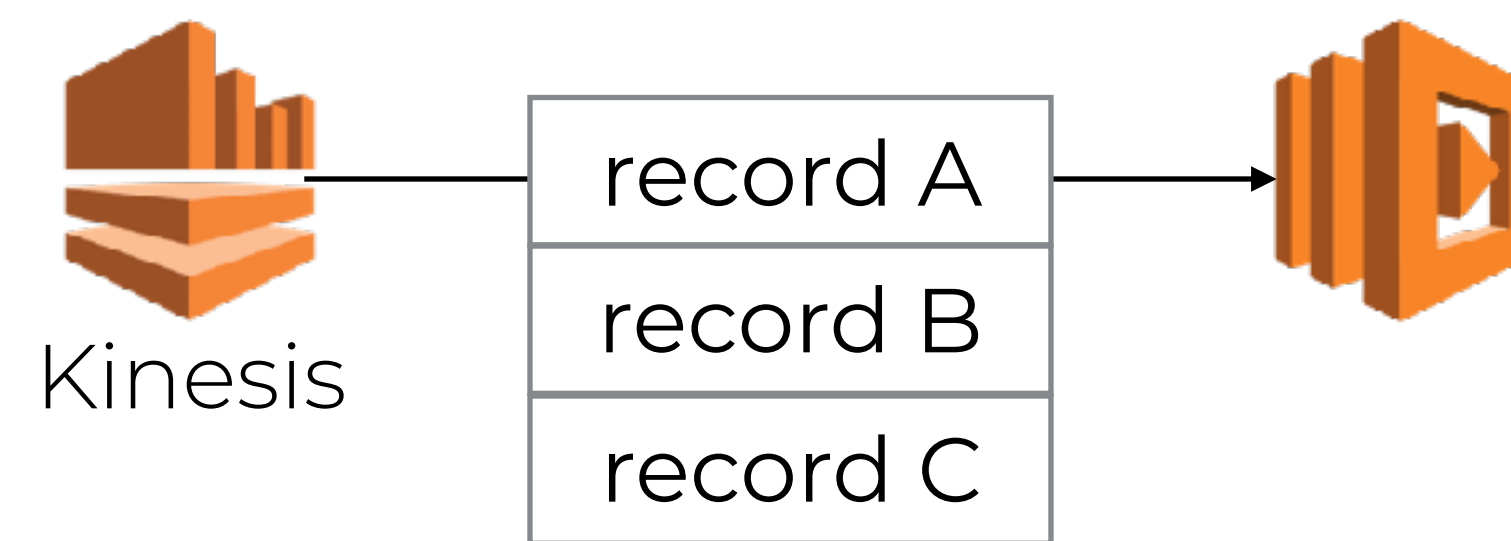
events should be processed in realtime  
(i.e. within a few seconds)

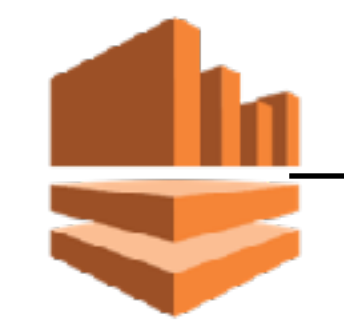
failed events should be retried, but the retries  
**should not violate** the realtime constraint



unprocessed events should be retrievable  
(i.e. available for human intervention, or root cause analysis)

- **Stream-based event sources** – For stream-based event sources (Amazon Kinesis Data Streams and DynamoDB streams), AWS Lambda polls your stream and invokes your Lambda function. Therefore, if a Lambda function fails, AWS Lambda attempts to process the erring batch of records until the time the data expires, which can be up to seven days for Amazon Kinesis Data Streams. The exception is treated as blocking, and AWS Lambda will not read any new records from the stream until the failed batch of records either expires or processed successfully. This ensures that AWS Lambda processes the stream events in order.



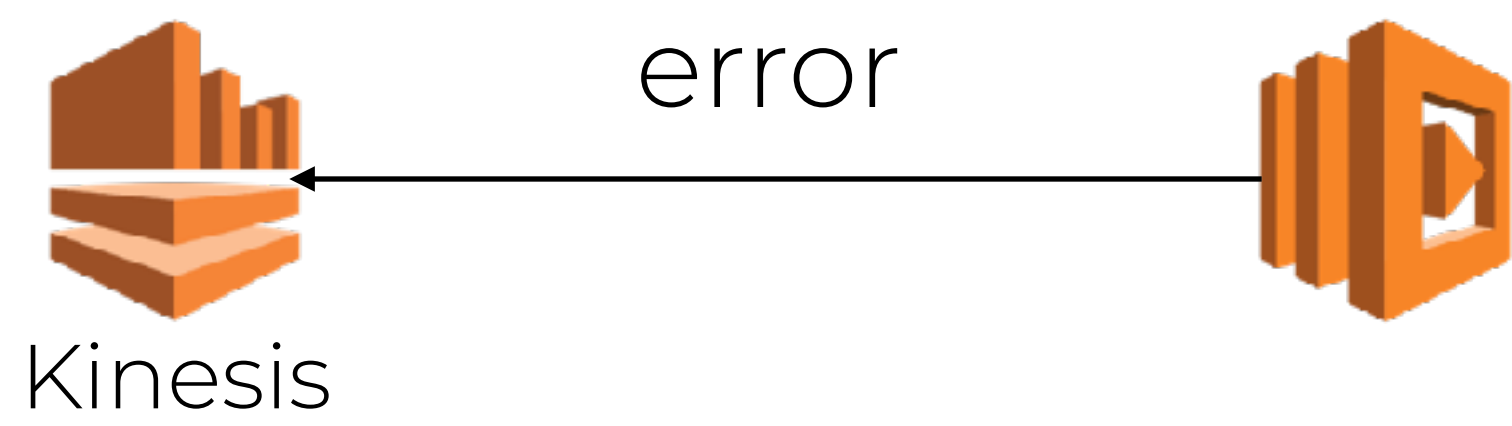


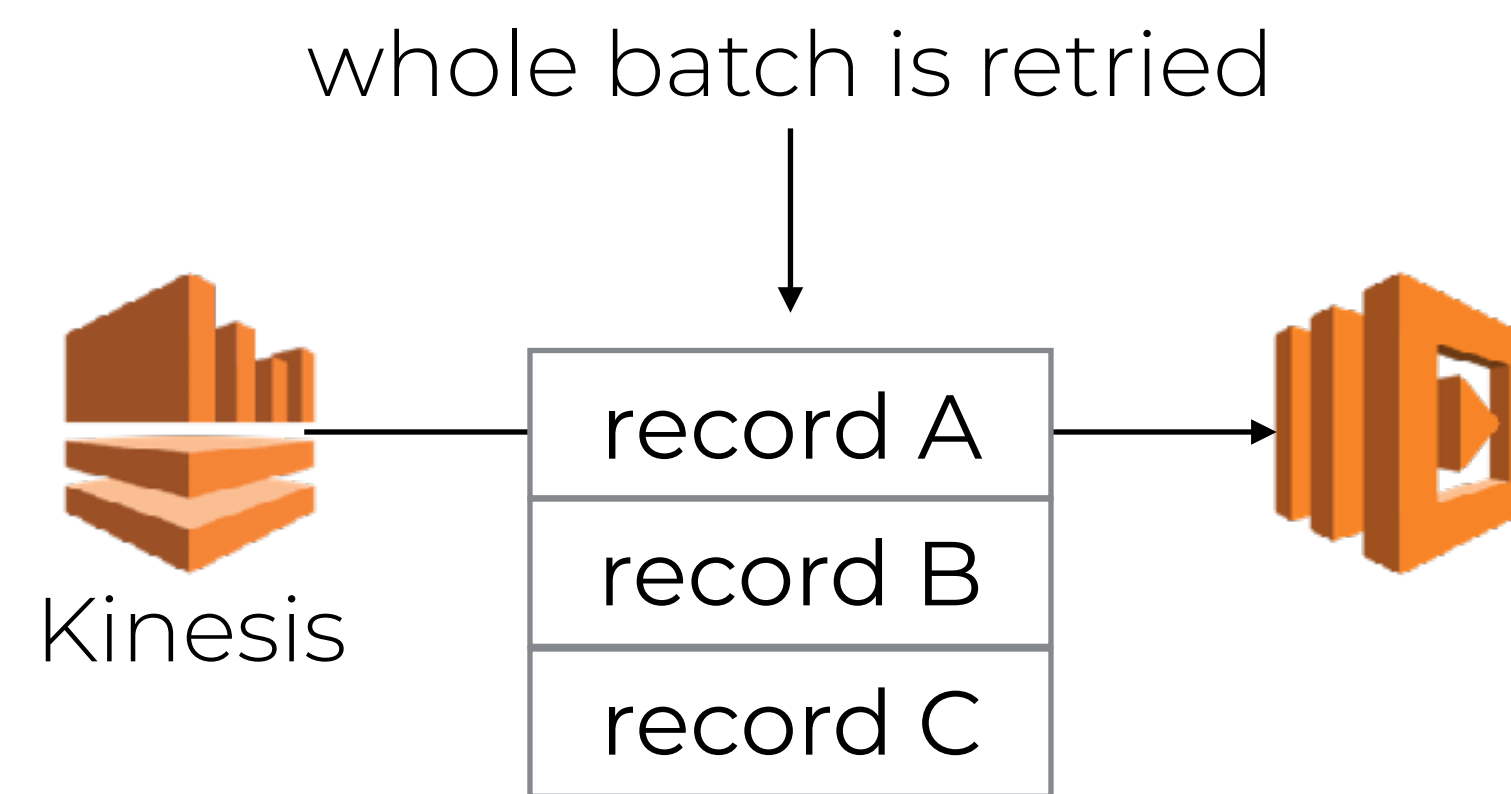
Kinesis



record A
record B
record C

← erred





you need to consider **Partial Failures** and **Idempotency** when processing Kinesis and DynamoDB streams with Lambda

failed events should be retried, but the retries  
**should not violate** the realtime constraint



failed events should be retried, but the retries  
**should not violate** the realtime constraint

you need to consider **Partial Failures** and **Idempotency** when processing Kinesis and DynamoDB streams with Lambda

**On-failure destination**

Lambda discards records that are expired or fail all retry attempts. You can send discarded records from a stream to an Amazon SQS queue or an Amazon SNS topic.

Queue or topic ARN

**Retry attempts**

The maximum number of times to retry when the function returns an error.

10000

**Maximum age of record**

The maximum age of a record that Lambda sends to a function for processing. The age can be up to 604,800 seconds (7 days).

604800

**Split batch on error**

If the function returns an error, split the batch into two and retry.

☐

**Concurrent batches per shard**

Process multiple batches from the same shard concurrently.

1

### On-failure destination

Lambda discards records that are expired or fail all retry attempts. You can send discarded records from a stream to an Amazon SQS queue or an Amazon SNS topic.

*Queue or topic ARN*

### Retry attempts

The maximum number of times to retry when the function returns an error.

10000

### Maximum age of record

The maximum age of a record that Lambda sends to a function for processing. The age can be up to 604,800 seconds (7 days).

604800

### Split batch on error

If the function returns an error, split the batch into two and retry.

☐

### Concurrent batches per shard

Process multiple batches from the same shard concurrently.

1

### On-failure destination

Lambda discards records that are expired or fail all retry attempts. You can send discarded records from a stream to an Amazon SQS queue or an Amazon SNS topic.

*Queue or topic ARN*

### Retry attempts

The maximum number of times to retry when the function returns an error.

10000

### Maximum age of record

The maximum age of a record that Lambda sends to a function for processing. The age can be up to 604,800 seconds (7 days).

604800

### Split batch on error

If the function returns an error, split the batch into two and retry.

☐

### Concurrent batches per shard

Process multiple batches from the same shard concurrently.

1

**On-failure destination**

Lambda discards records that are expired or fail all retry attempts. You can send discarded records from a stream to an Amazon SQS queue or an Amazon SNS topic.

Queue or topic ARN

**Retry attempts**

The maximum number of times to retry when the function returns an error.

10000

**Maximum age of record**

The maximum age of a record that Lambda sends to a function for processing. The age can be up to 604,800 seconds (7 days).

604800

**Split batch on error**

If the function returns an error, split the batch into two and retry.

☐

**Concurrent batches per shard**

Process multiple batches from the same shard concurrently.

1

6, 3, 1, 1, 1, 1, ...

On-failure destination

Lambda discards records that are expired or fail all retry attempts. You can send discarded records from a stream to an Amazon SQS queue or an Amazon SNS topic.

Queue or topic ARN

only count the "same" batch

Retry attempts

The maximum number of times to retry when the function returns an error.

10000

Maximum age of record

The maximum age of a record that Lambda sends to a function for processing. The age can be up to 604,800 seconds (7 days).

604800

6, 3, 1, 1, 1, 1, ...

Split batch on error

If the function returns an error, split the batch into two and retry.

☐

Concurrent batches per shard

Process multiple batches from the same shard concurrently.

1



```
{
  "Body": "{
    \"requestContext\":{
      \"requestId\": \"0cdf90e8-6290-4834-a46c-cffdd6277c4f\",
      \"functionArn\": \"arn:aws:lambda:us-east-1:374852340823:function:aws-sdk-example\",
      \"condition\": \"RetryAttemptsExhausted\",
      \"approximateInvokeCount\": 11
    },
    \"responseContext\":{
      \"statusCode\": 200,
      \"executedVersion\":
        \"$LATEST\",
      \"functionError\":
        \"Unhandled\"
    },
    \"version\": \"1.0\",
    \"timestamp\": \"2020-02-08T19:53:20.998Z\",
    \"KinesisBatchInfo\":{
      \"shardId\": \"shardId-000000000000\",
      \"startSequenceNumber\": \"49604045060801379268037077355262289118696897411062693890\",
      \"endSequenceNumber\": \"49604045060801379268037077355262289118696897411062693890\",
      \"approximateArrivalOfFirstRecord\": \"2020-02-08T19:47:05.582Z\",
      \"approximateArrivalOfLastRecord\": \"2020-02-08T19:47:05.582Z\",
      \"batchSize\": 1,
      \"streamArn\": \"arn:aws:kinesis:us-east-1:374852340823:stream/yc-test\"
    }
  }"
}
```



```
{
  "Body": "{
    \"requestContext\":{
      \"requestId\": \"0cdf90e8-6290-4834-a46c-cffdd6277c4f\",
      \"functionArn\": \"arn:aws:lambda:us-east-1:374852340823:function:aws-sdk-example\",
      \"condition\": \"RetryAttemptsExhausted\",
      \"approximateInvokeCount\": 11
    },
    \"responseContext\":{
      \"statusCode\": 200,
      \"executedVersion\":
        \"$LATEST\",
      \"functionError\":
        \"Unhandled\"
    },
    \"version\": \"1.0\",
    \"timestamp\": \"2020-02-08T19:53:20.998Z\",
    \"KinesisBatchInfo\":{
      \"shardId\": \"shardId-000000000000\",
      \"startSequenceNumber\": \"49604045060801379268037077355262289118696897411062693890\",
      \"endSequenceNumber\": \"49604045060801379268037077355262289118696897411062693890\",
      \"approximateArrivalOfFirstRecord\": \"2020-02-08T19:47:05.582Z\",
      \"approximateArrivalOfLastRecord\": \"2020-02-08T19:47:05.582Z\",
      \"batchSize\": 1,
      \"streamArn\": \"arn:aws:kinesis:us-east-1:374852340823:stream/yc-test\"
    }
  }"
}
```

#### On-failure destination

Lambda discards records that are expired or fail all retry attempts. You can send discarded records from a stream to an Amazon SQS queue or an Amazon SNS topic.

*Queue or topic ARN*

#### Retry attempts

The maximum number of times to retry when the function returns an error.

10000

#### Maximum age of record

The maximum age of a record that Lambda sends to a function for processing. The age can be up to 604,800 seconds (7 days).

604800

#### Split batch on error

If the function returns an error, split the batch into two and retry.

☐

#### Concurrent batches per shard

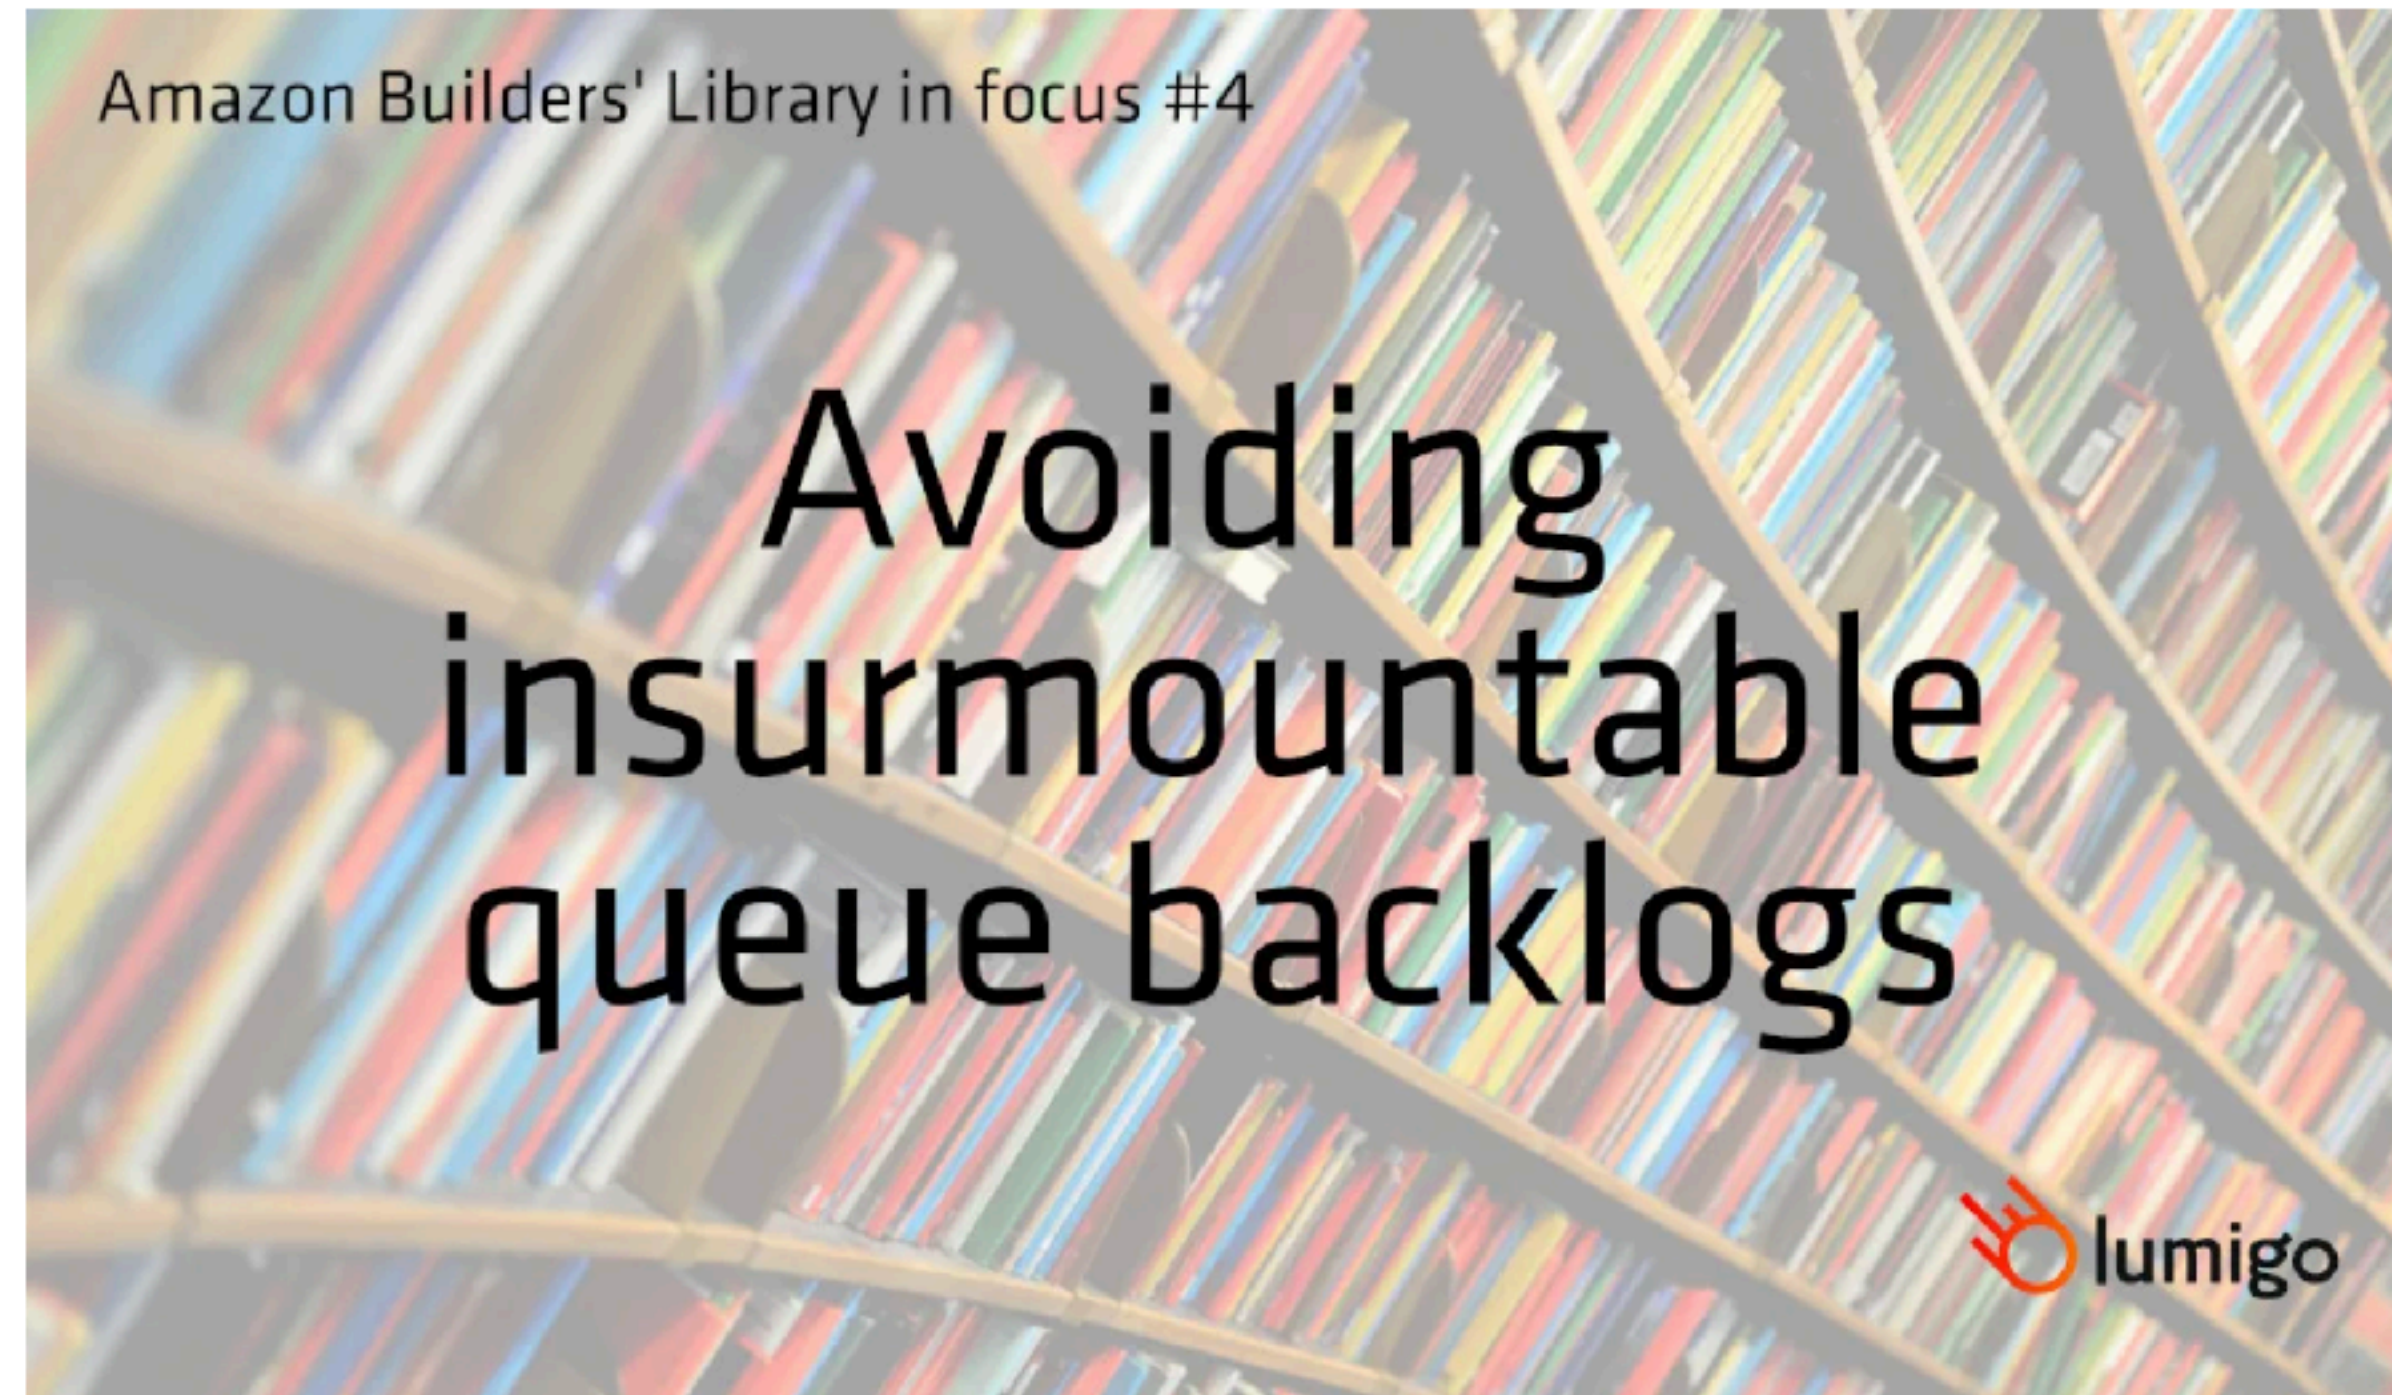
Process multiple batches from the same shard concurrently.

1



## AMAZON BUILDERS' LIBRARY IN FOCUS #4: AVOIDING INSURMOUNTABLE QUEUE BACKLOGS

---



In the latest article in our series focusing on the **Amazon Builders' Library**, Yan Cui highlights the key insights from **Avoiding insurmountable queue backlogs** by AWS Principal Engineer David Yanacek.

<https://lumigo.io/blog/amazon-builders-library-in-focus-4-avoiding-insurmountable-queue-backlogs>



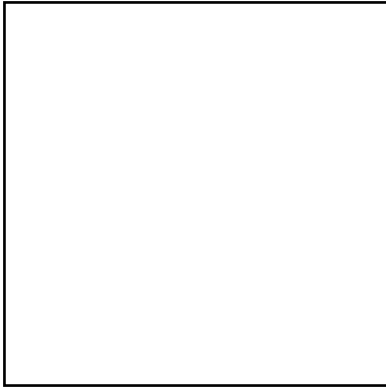
**SQS**



**Lambda**



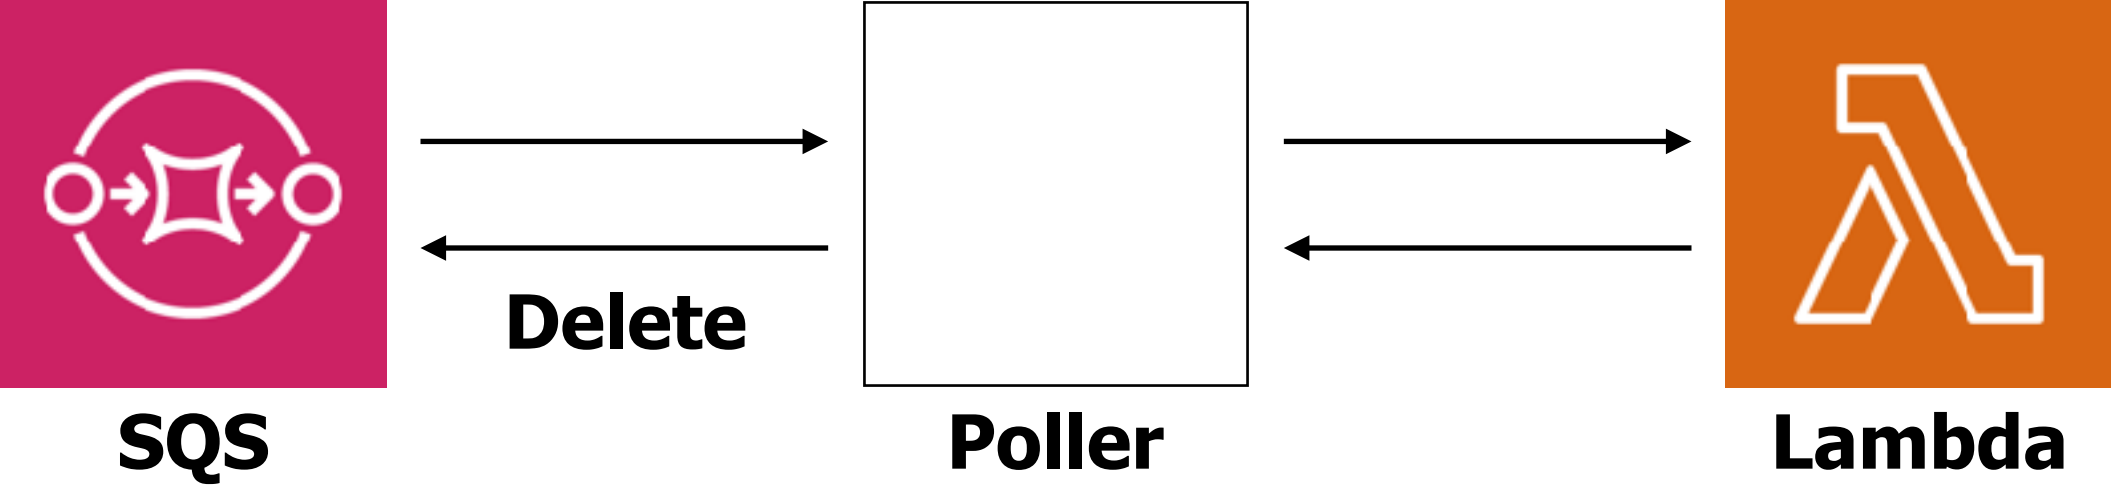
**SQS**

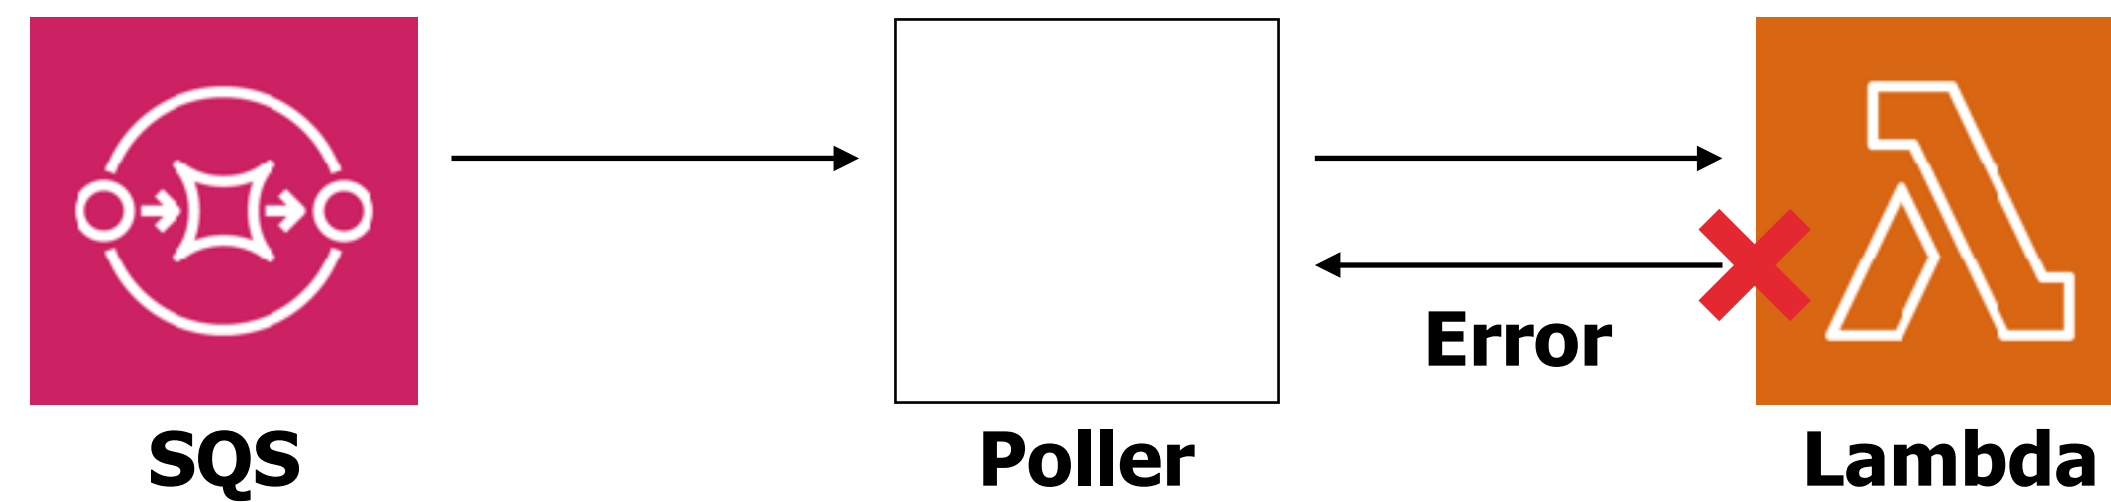


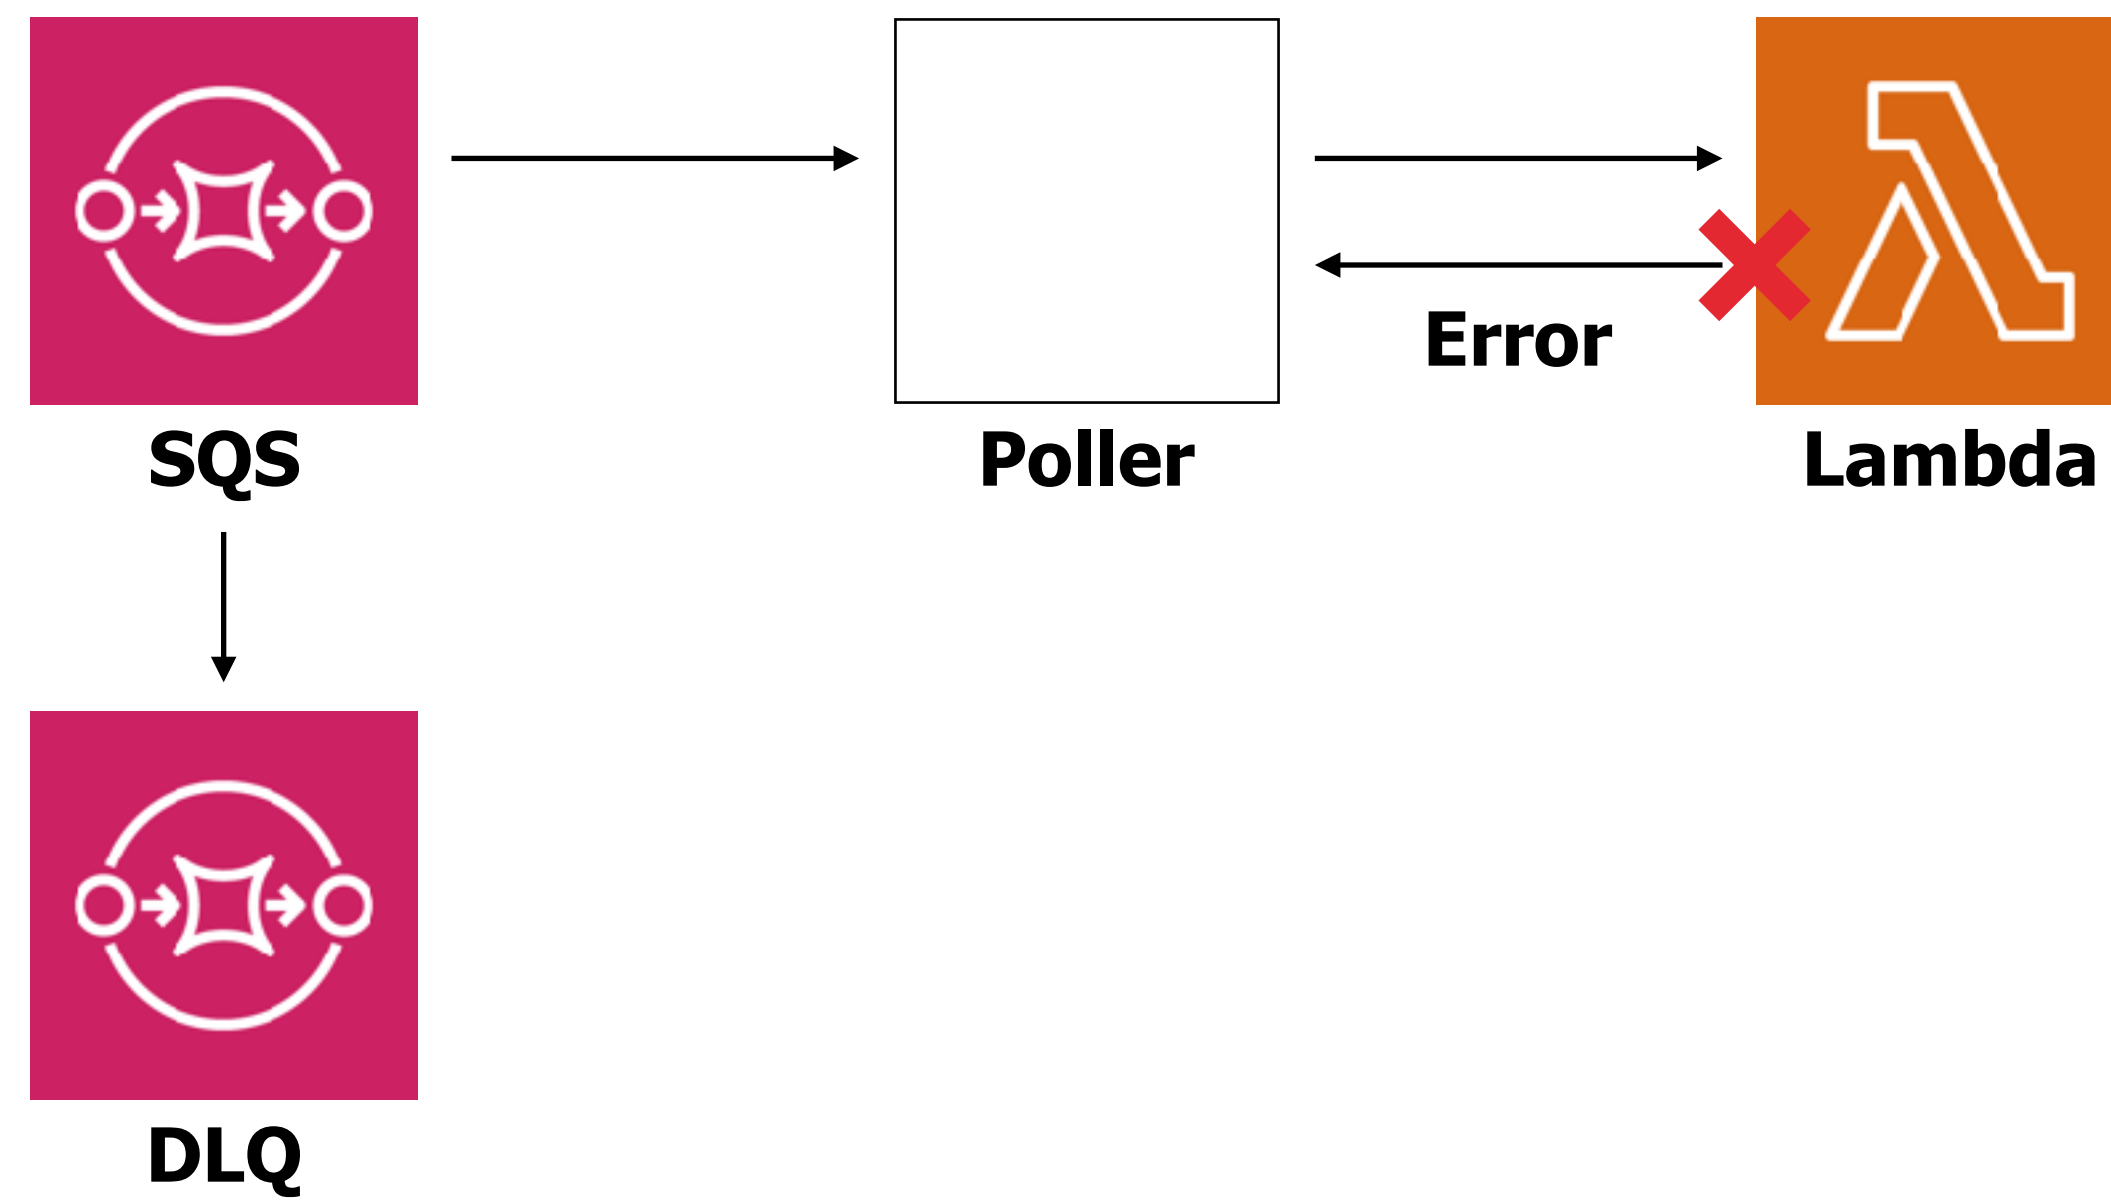
**Poller**



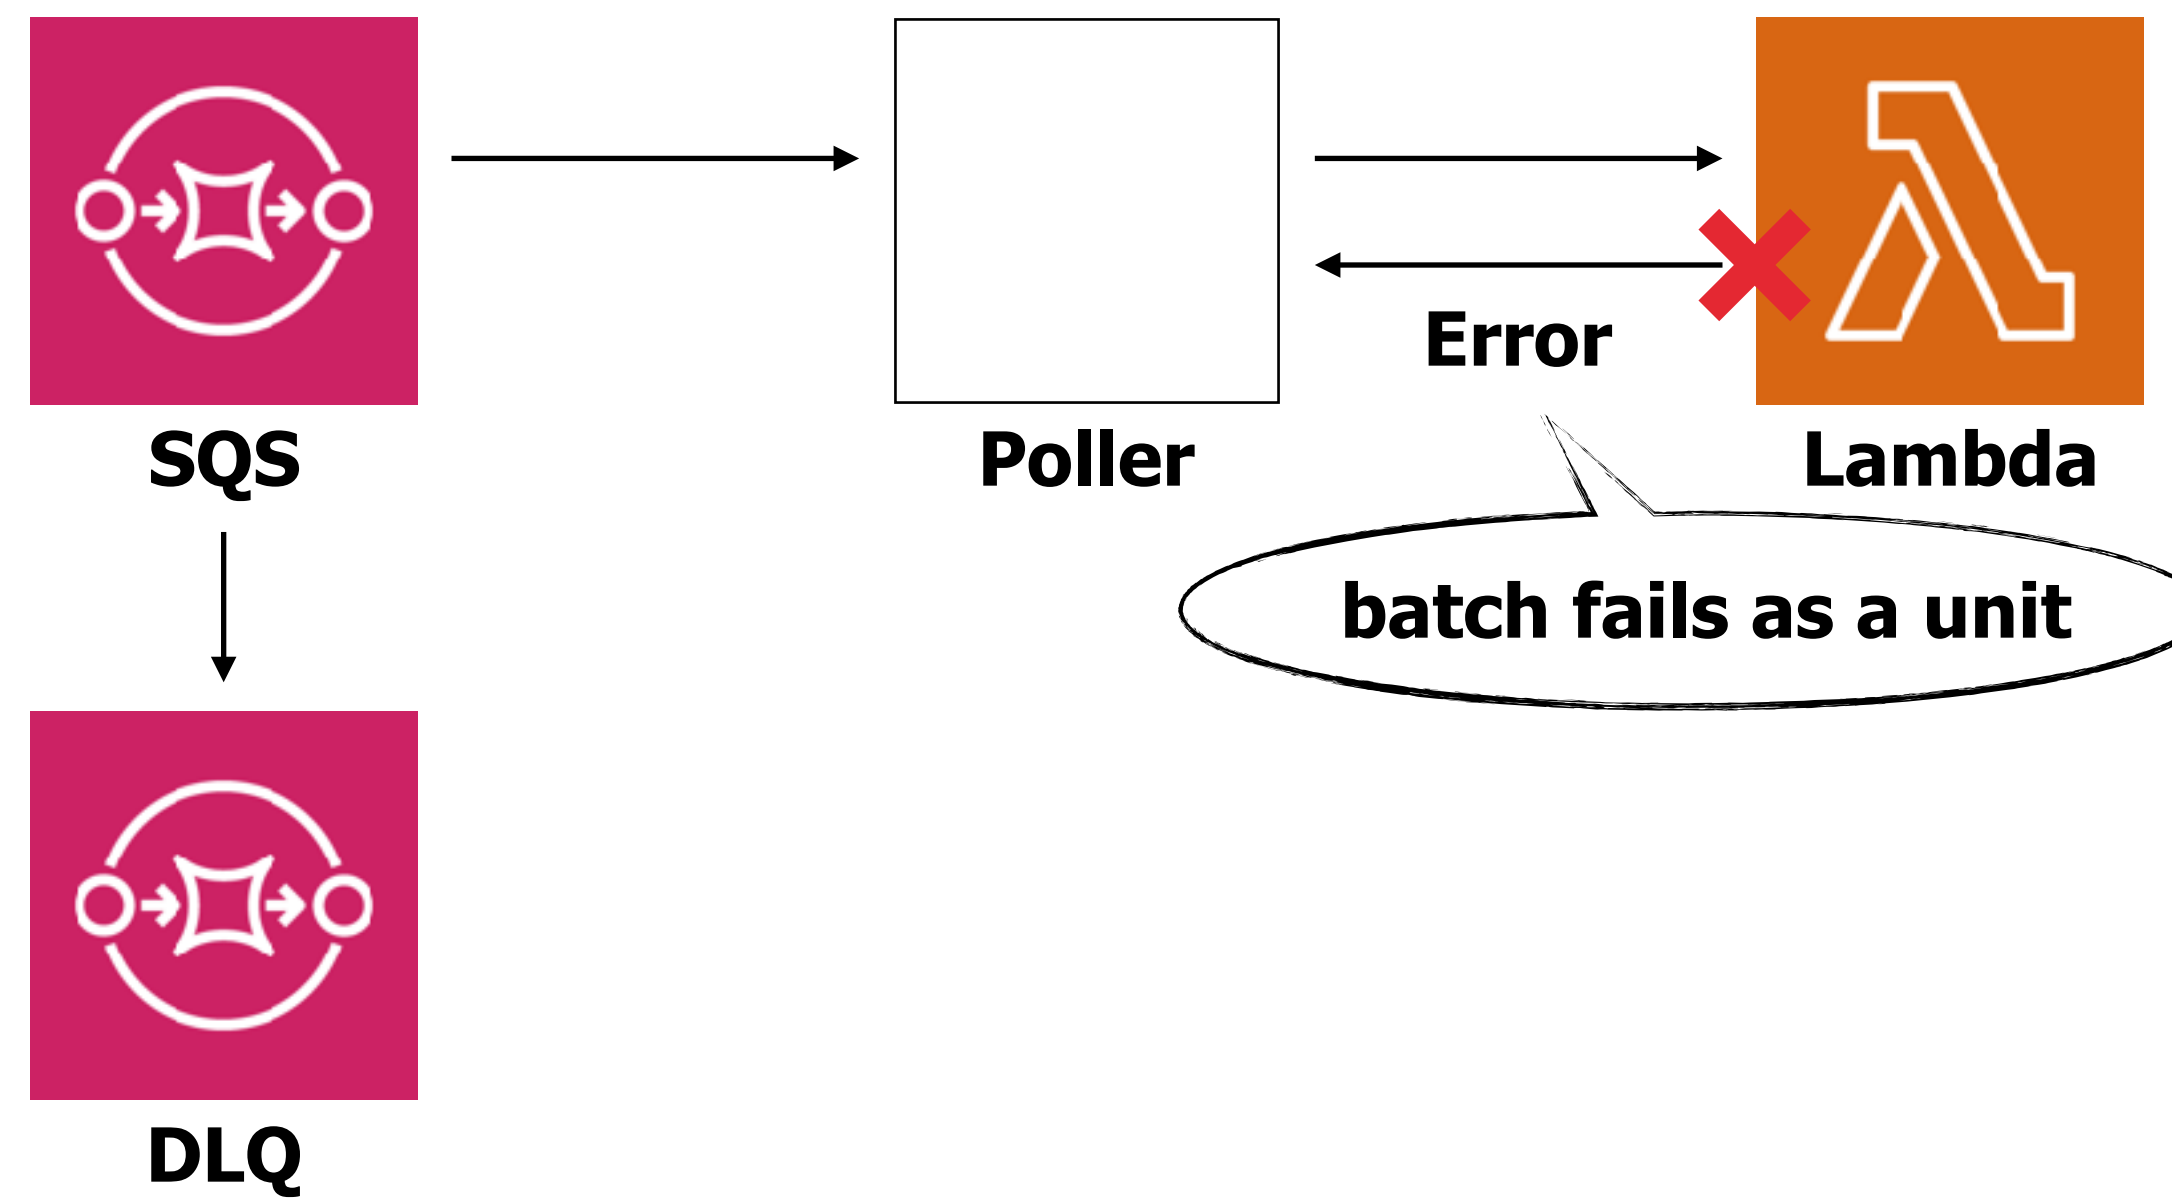
**Lambda**



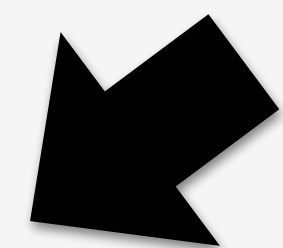








```
{
  "Type" : "AWS::Lambda::EventSourceMapping",
  "Properties" : {
    "BatchSize" : Integer,
    "BisectBatchOnFunctionError" : Boolean,
    "DestinationConfig" : DestinationConfig,
    "Enabled" : Boolean,
    "EventSourceArn" : String,
    "FilterCriteria" : Json,
    "FunctionName" : String,
    "FunctionResponseTypes" : [ String, ... ],
    "MaximumBatchingWindowInSeconds" : Integer,
    "MaximumRecordAgeInSeconds" : Integer,
    "MaximumRetryAttempts" : Integer,
    "ParallelizationFactor" : Integer,
    "Queues" : [ String, ... ],
    "SelfManagedEventSource" : SelfManagedEventSource,
    "SourceAccessConfigurations" : [ SourceAccessConfiguration, ... ],
    "StartingPosition" : String,
    "StartingPositionTimestamp" : Double,
    "Topics" : [ String, ... ],
    "TumblingWindowInSeconds" : Integer
  }
}
```



ReportBatchItemFailures

```
{
  "batchItemFailures": [
    {
      "itemIdentifier": "id2"
    },
    {
      "itemIdentifier": "id4"
    }
  ]
}
```